
Title of the project

Subtitle

Project Report
Group: SW805F20

Aalborg University
Department of Computer Science
Selma Lagerlöfs Vej 300
9220 Aalborg East, DK

Copyright © Aalborg University 2019

Photographic, mechanic or any other form of duplication of this paper is not allowed according to Danish copyright law and without permission by the authors.

Department of Computer Science

Aalborg University
Selma Lagerlöfs Vej 300
9220 Aalborg East, DK
www.cs.aau.dk

Title:

Title of the project

Abstract:

This is the best abstract ever written
--

Theme:

Mobility

Project Period:

Spring Semester 2020

Project Group:

SW805F20

Participant(s):

Andreas Stenshøj
Daniel Moesgaard Andersen
Frederik Valdemar Schrøder
Jens Petur Tróndarson
Rasmus Bundgaard Eduardsen
Mathias Møller Lybech

Supervisor(s):

Brian Nielsen

Copies: 1

Page Numbers: 67

Date of Completion:

May 28th, 2020

The content of this report is freely available, but publication (with reference) may only be pursued due to agreement with the author.

Contents

1	Introduction	1
1.1	Project idea	1
1.2	Essence	2
2	Sprint 1	5
2.1	Sprint goal and introduction	5
2.2	Our process	5
2.3	Prototypes	7
2.4	Architecture of the game	9
2.5	Pozyx	10
2.6	Unity introduction	14
2.7	Networking	14
2.8	Experiment with Pozyx	25
2.9	Sprint 1 conclusion	32
3	Sprint 2	35
3.1	Sprint goal and introduction	35
3.2	Deployment Diagram	35
3.3	Accessing games on the network	37
3.4	UDP implementation	38
3.5	Current product	40
3.6	Retrospective on the process	41
4	Sprint 3	45
4.1	Sprint 3 goals	45
4.2	Dead reckoning	45
5	Sprint 4	49
6	Appendix	51
6.1	Results for experiment	51
6.2	Experiment with five tags	55
6.3	Test plan	58

Bibliography	61
List of Figures	63
List of Tables	65

Terms and abbreviations

Terms and abbreviations used in the report:

Pozyx : The hardware used for positioning.

AR : Augmented reality.

UWB : Ultra-wide bandwidth.

TWR : Two-way-ranging.

OSI : Open systems interconnection.

GDOP : Geometric Dilution of Precision

TCP : Transmission Control Protocol.

UDP : User Datagram Protocol.

Chapter 1

Introduction

1.1 Project idea

The idea for this project is to create a location-based competitive game using augmented reality (AR). Two teams will compete against each other to score the most goals using a ball. Each player will be equipped with a smartphone-based virtual reality headset, and these will display the playing field from a top-down 2D view. To achieve this, each player's position needs to be tracked as well as where the ball is located on the field. In the top-down view, each player needs to see the positions of the other players and the ball. They also need to see their position on the playing field and where the goals are. The players should be able to set a number of goals they need to score to win before beginning the game.



Figure 1.1: An illustration of the playing field

In Figure 1.1, an illustration of the playing field for the game is shown. There are goals at each end of the field, and the teams score goals by getting the ball between the goalposts. An alternative version of the game is suggested where, instead of goals at each end of the field, there would be virtual goal zones seen in the game which the teams need to bring the ball into. These zones could even change locations as the game progressed.

1.1.1 Technical requirements for the project

Multiple pieces of hardware will be required to realize the vision of the game. First of all, each player must have a headset to hold their smartphone such that they can view a virtualized version of the playing field while playing. In order for the data to be synchronized between the players, they will also need to be equipped with a positioning device, which can transmit their location to the other players. This transfer of information will require a networking solution so that the virtualized playing field is synchronized between the players.

1.1.2 Problems to consider

The project idea proposes some problems that will need to be solved for the game to work. We will need to consider which technologies to use for the development of the visual aspect of the game which should show a top-down view for each player. As it is something we do not have experience with, it would be preferable if it is not necessary to have to build it from scratch. We will also need hardware that can track the positions of players and the ball. This must be accurate and update quickly such that the players do not run into each other, otherwise the game will not work. Another problem to consider is how the ball should be displayed in a 2D view. For the players to be able to find the ball on the field, it either has to be quite large to make it easier to find from the top-down view, or the game will need some metric to display how far the ball is from the ground. The game will also need to track when the ball has crossed the goal line and then give feedback to the players. Another problem to solve is how to keep the positional data synchronized across all the players' devices, as it will be difficult to play the game without accurate data.

1.2 Essence

For the process of project development, we have chosen to work with Essence. Previous semesters we have worked with an agile approach inspired by Scrum, however, this semester we are attempting to apply the Essence approach. The basic idea of Essence is to encourage diverse thinking in the team, even though all members of our team share a similar background as bachelors in Software.

Essence uses two strategies to support value creation:

- *A systematic use of diverse viewpoints.* Values, views, and roles are used frequently in Essence. By using different views and roles to represent problems and solutions, Essence tries to facilitate a range of viewpoints on how a problem needs to be understood and solved.
- *A focus on idea maturation more than idea generation.* Essence applies the concept that ideas develop over time and tries to stimulate the team to evaluate and refine ideas [1].

1.2.1 Four variants of innovation

Essence tries to support innovation in software development, and hereby it defines four different variants of innovation, which are: [1]:

- *Product innovation* is new or radically changed software products or services.
- *Process innovation* is software solutions offering the user new or radically improved ways to produce products or services.
- *Project innovation* is fitting software solutions from earlier projects into new application domains
- *Paradigm innovation* is about software solutions coupled with changes in the mental model of what a business is, who the users are or what the market is.

1.2.2 Paradigms

There are two well-established software development paradigms: the document-oriented paradigm, which we know from the waterfall approach and agile paradigm which we know from for example extreme programming and Scrum. The author of Essence considers the new emerging paradigm called the pragmatic paradigm.

The document-oriented paradigm portrayed software developers as being document-oriented. The requirements are static and allow for a top-down waterfall approach to software development, which pays small attention to creativity and innovation.

The agile paradigm sees software development as user-oriented. Requirements are presumed dynamic as customers learn about options and constraints within the course of the project. This leads to incremental software development.

The pragmatic paradigm is problem-oriented. Systems are becoming more complex and it is more difficult to separate systems from each other. The amount of data, software libraries and hardware components available is steadily increasing, leading to hypercomplex software projects. Hypercomplexity is a degree of complexity that makes it impossible to make rational decisions within a reasonable time constraint. The most important features of this paradigm are that requirements are not completely known when the projects start. Ideas evolve in the process of the project, and during the project the requirements for the project are negotiable [1].

1.2.3 Core concerns

All software projects involve these four core concerns:

- Do we know what to build?
- Do we understand the solution?

- Do we understand the problem?
- Should we pivot or persevere?

1.2.4 Team organization

Within the team organization, in Essence, roles are used to create heterogeneity in teams, to ensure diverse points on views and to ensure cohesion despite diversity. The focus of these roles is to increase learning with personal interaction by sharing insights and experiences. The roles also ensure that the team understands the problem domain, and see potentials in the technology domain.

As a rule of thumb, the roles are persistent meaning that a member will have the same role for the duration of the project. The roles in Essence are compatible with agile software development, making it possible to combine Essence with other processes like Scrum.

There are four roles in essence:

- Child
- Responder
- Challenger
- Anchor

The role of *Child* can ask any questions and make propositions that are opposite of previous decisions. The rest of the team is not allowed to criticize the child, but they are however allowed to ignore the person's suggestions. The child is one of the main sources of ideas and other perspectives on the project. Outsiders are also allowed to take this role.

Responders are the developers in the team, and they are usually the majority within the team. Responders work closely together with the *Challenger*, so that the most important features are developed first.

Challenger is the customer or customer representative. The challenger can be compared to the *Product Owner* in Scrum. This role formulates and explains the Challenge, prioritizes features and accepts the solutions. There can be more than one Challenger, but if there are they must agree on the product vision.

Anchor is the one responsible for leading evaluations but does not decide the consequences. If necessary, the anchor can intervene and remove threats to the team's ability to develop ambitious responses. A potential threat could be something that results in productivity issues.

Chapter 2

Sprint 1

2.1 Sprint goal and introduction

The goal for this sprint is to explore the project idea and look for possible solutions to some of the problems that the project idea introduces. In this sprint, we want to learn more about visualizing the game for the users so they can see themselves in the game as well as the playing field.

We will look into using Pozyx to define the corners of the playing field as well as get the players' positions in an accurate way with fast updates, and to determine whether or not Pozyx is even viable for use in the context of this project an experiment will be conducted where we look into the accuracy of the positional precision of the tags. As this is the first sprint of the project, the main focus is to gain more knowledge about the various aspects of the project and to gain a shared vision amongst the group members as to how the game should work. This initial shared vision will be achieved by making prototypes of the game as well as architectural diagrams of how different components in the game will work together.

2.2 Our process

As described in [section 1.2](#), a secondary focus of this project will be to attempt using the Essence process model as taught by the Software Innovation course. We have been working with a ScrumBut approach in previous semesters and at the beginning of this semester, but we decided that we wanted to try a different process this semester, and as we learned more about the Essence approach we decided to pivot, changing our process.

However, we have chosen to keep sprints, stand-up meetings and retrospective meetings from Scrum as these can complement Essence. To make the report fit this format, it has been split into 4 sprints that we will go through during the semester, where each sprint has a length of 3 weeks.

2.2.1 Roles

We have divided the roles described in subsection 1.2.4 between the members of the group. One member has the *Challenger* role, meaning that they are accountable for prioritizing the tasks in the backlog. The process of prioritizing tasks is further described in the following section. Another member has the *Anchor* role, and is responsible for changes to the process and is in charge of leading evaluations of the process. The rest of the group functions as *Responders*, which is the role for the developers of the project. The *Child* role fluctuates between members of the group. Everyone can add suggestions to improve an idea and give other perspectives on the project.

Due to our team size, the challenger and anchor will also work as developers during the duration of the project.

2.2.2 Prioritizing tasks

Our backlog is saved as a board on Jira, which can be seen on Figure 2.1. The leftmost column is the *Suggested* column. Everyone can make suggestions for tasks that they find useful for the project. After each stand-up meeting, the challenger will present new suggestions that seem relevant to work on in the near future. This presentation will include the definition of done, and all members of the group will then vote on how valuable it is for the project and how time-consuming it is. The priority of a task is then calculated as $\text{reward} - \text{time}$, which is an arbitrary number to indicate how important it is.

The challenger then chooses the most important features from the *Discussed* column, often based on the highest priority, as *Chosen for Development*. Responders then have the opportunity to take tasks from this column and move it to the next column *In progress* when they start working on it. When the task has been completed, reviewed and merged into the develop branch, it is automatically moved to the column *Done*.



Figure 2.1: The board for tasks.

2.2.3 Pair reviews

To ensure quality and make sure that people follow the test plan, everything undergoes a formal review before it is added to the project. The test plan is our description and an agreement of what needs to be tested and how thoroughly it should be tested. More details about the test plan can be found in [section 6.3](#). For report tasks, two reviewers are assigned to review and approve it before it can be merged. For coding tasks, two people are likewise assigned to review it, but the review has to be done as a pair, meaning that they have to physically sit together and go through the code on a shared screen. These pair reviews are a good way to share knowledge about the implementation through the group, as people have to understand it to be able to discuss it. In previous semesters people have reviewed the code separately, which leads to fewer comments as the code was not discussed between reviewers.

2.3 Prototypes

The first step in the project is to arrive at a common vision about the layout and functionality of the system. To accomplish this, a series of wireframe prototypes were created. A wireframe is a lo-fi prototyping technique where you create a grey box schematic that outlines the primary features of the user interface, unlike hi-fi prototypes which generally have a higher amount of detail and are used to test the interaction with the application. The primary focus of these prototypes is not to be used for implementation, but rather for comparing opinions about how the flow in the system should be created. Since the user interface is mostly focused on the mobile devices that the players will be wearing, it was decided that the host computer

should simply have a text-based interface, as seen on Figure 2.2.

```

Terminal
> Amount of players:
4

> Position of anchor 1 (mm)
0 0 600

> Position of anchor 2 (mm)
0 10000 1800

> Position of anchor 3 (mm)
10000 0 1800

> Position of anchor 4 (mm)
10000 10000 1800

> Ball tag ID
0x86ef

> Player 1 tag
0x6492

> Player 2 tag
0x6421

> Player 3 tag
0x6195

> Player 4 tag
0x9fe2

> Confirm selection? (Y / N)
Y

Waiting for player connections... (IP: 192.168.153)

Player 1 / 4 connected (Given tag: 0x6492, Team: Blue)
Player 2 / 4 connected (Given tag: 0x6421, Team: Red)
Player 3 / 4 connected (Given tag 0x6195, Team: Blue)
Player 4 / 4 connected (Given tag: 0x9fe2, Team: Red)

> All players connected, start game? (Y / N)
Y

```

Figure 2.2: Prototype of hosting interface

When a user starts the application, they will be greeted with an input field, where they will specify the IP address of the host machine, as seen on Figure 2.3.



Figure 2.3: Prototype of game menu

After inputting an IP and confirming, they will be redirected to a page where they can see how many users have connected to the host, as seen on Figure 2.4.



Figure 2.4: Prototype of screen where a user has connected to the host

When all users have connected, the host can start the game, and they will now see the virtual game field, as seen on Figure 2.5. In this prototype, the players' icons are represented by small squares and the goals are represented by larger rectangles and have the same color as the players they belong to.

The current player's icon is highlighted by having a solid color, whereas the other players are just shown as outlines. The players are divided into two teams which are represented by the color of the player icons. In the middle of the screen is the ball in a designated starting area to make the game fair for both teams.




Figure 2.5: Prototype of in-game screen

2.4 Architecture of the game

In addition to the prototypes, which served as a broad overview of how the flow in the application should be, we will look at the overall architecture of the different components involved in the game from a more technical perspective.



Figure 2.6: The different components of the game

In Figure 2.6 we see an illustration of the initial idea of how the different components of the game are going to work together. The arrows in the diagram show the flow of data in the game. The Pozyx component is responsible for providing the positional data to the host, this will be described in further depth in section 2.5. Each client/-player is equipped with one of the tags that track their positions on the field. The playing field’s dimensions will be set to match the distances between the anchors. The host has the positional data for each client, but it will also need to know which of the tags each client is equipped with, such that it can continuously transfer positional data to the clients. This information will also include which of the tag ids belong to which clients. Each client will be running an instance of the game, and with the use of the positional data, it will render the playing field, the ball and each of the players. The host is responsible for checking if the ball has crossed the goal line and it should then provide this information to the players. In general, the clients’ instances of the game should only be responsible for rendering.

Any game logic should happen on the server-side to make sure that the game is synchronized for each player.

2.5 Pozyx

Since positioning is a major part of the project, it is important to have accurate positioning information. Pozyx is a hardware/software solution that is used to provide positioning with an accuracy of down to 10 cm [7]. It makes use of ultra-wide band (UWB) in combination with machine learning for positioning, which according to

their documentation is more precise and efficient than traditional positioning systems such as WiFi, Bluetooth, RFID, and GPS.

Since the two major requirements for positioning in this project are precision and a high update rate to ensure that the players can have reliable data available, the Pozyx system seems like a good place to start.

The Pozyx tags support update rates of up to 125 Hz for a single tag [7]. The Creator system from Pozyx is sold with 4 anchors and 5 locatable tags. An anchor is a stationary sensor used by the moveable tags to get their exact position.

2.5.1 Finding the location of anchors

A trilateration method is used for finding the position of a given tag using the anchors. This method uses basic geometry to estimate the position by measuring the distance to the anchors of which we know the position. With this distance estimate, it is possible to draw a circle with a given radius. If we use two anchors, we will have two intersection points which are the possible positions of the tag. This means that to find a two-dimensional location, we will need at least three anchors, which will lead to only a single point where all three circles intersect, as seen on [Figure 2.7](#).



Figure 2.7: An example of trilateration using three anchors

The issue with this approach is that the measurements are not perfect, which might cause the circles to not intersect at exactly one point which will make the positional data seem to jitter.

2.5.2 Using UWB

To find the position of the tags, Pozyx makes use of radio waves. Radio waves travel at the speed of light, so by dividing the time of travel between anchors with the speed of light, the distance between them can be found.

Because the speed of light is so fast the time measure needs to be very accurate to get the correct distance. To achieve this the anchors make use of UWB [9]. UWB is a technology for transmitting data over a wide bandwidth usually wider than 500 MHz. The Pozyx sensors use exactly 500 MHz which means they just qualify as a UWB system. Because the bandwidth is this wide it makes the wavelength very short, and by combining multiple sinusoidal signals with slightly different frequencies the Pozyx tags can create a pulse with a peak which is very narrow. Measuring a narrow peak results in a more accurate timing, allowing the Pozyx to be much more precise than other technologies like Bluetooth and WiFi. High bandwidth means faster data transfer, which most people would prefer, but if everyone were to use the same frequency the signals would interfere with each other, therefore the use of high-frequency signals is tightly regulated[8]. To reduce the amount of interference the regulations say that UWB systems may only transmit at very low power. This means that a UWB transmission can not travel very far and therefore the chance of it interfering with another system is very low. Because of this we are not able to build huge playing fields with the limited amount of anchors we have if we also want it to be precise or even playable. Building a playing field that is too big could cause some tags to not reach all anchors and therefore not calculate a position. This could be combated somewhat by having more anchors, but as will be explained in 2.5.4 the slave tags might not be able to communicate with the master tag because the distance between them is too large.

2.5.3 Alternatives to UWB

Having a high accuracy is essential for the users' experience since a low accuracy could result in the in-game information being incorrect or just too imprecise to make it a joyful experience. While Pozyx uses UWB for positioning, it is interesting to take a look at the primary alternatives to indoor positioning and to see if any of them can compete with Pozyx' high accuracy [7].

GPS

The obvious alternative for positioning data is GPS, which is frequently used for outdoors positioning. In optimal outdoor situations, GPS can provide accuracy within 4.9 meters [16]. However, the GPS positioning accuracy can be degraded due to buildings, bridges, and trees, making it a sub-optimal choice for indoor positioning.

Wi-Fi

Since the project already utilizes networking for data transfer, Wi-Fi may be worth considering to decrease the number of technologies needed. However, Wi-Fi only provides an accuracy of between 5 and 15 meters depending on the hardware chosen for access points and clients.

In addition to this, iOS devices are blocked from using Wi-Fi for indoor navigation purposes, since it usually makes use of a technology called fingerprinting, which only functions with Android devices due to technical restrictions [10].

Bluetooth

Using Bluetooth for indoor location is similar to how Pozyx works with anchors and tags. Instead of anchors, Bluetooth beacons are able to send out signals with a range of up to 30 meters [11]. With Bluetooth, you get a significantly higher accuracy compared to GPS and Wi-Fi. In the optimal settings, this solution can provide an accuracy of up to one meter. Unlike Wi-Fi, this solution would work on both iOS and Android devices.

RFID

Finally, we have RFID which uses radio waves to wirelessly transmit the identity of an object. Unlike the other solutions, RFID offers high accuracy, but a very limited range of less than a meter [12].

For a game like what we are building, it would be possible to use RFID for the location of the ball to check which player is currently holding it, but due to the limited range, it would not make sense to use it for localizing the players.

Conclusion

Generally, using Pozyx with their UWB technology for tracking the position of the players and the ball appears to be the most optimal solution.

2.5.4 Two-way-ranging

We are using the Pozyx Creator Kit Lite which uses the Two-way-ranging (TWR) protocol for positioning [17].

A tag calculates its position by communicating with the anchors one by one, getting the distance from the anchor to itself. Once it has the distance from 3 anchors it can compute its position utilizing trilateration.

If multiple tags are being used at once, one tag is made the master tag and the other tags become the slave tags. The master tag instructs the slave tags to report their position to the master tag one by one. The master tag is then usually connected to a computer that can use the position data. This technique does not scale well as all

the slave tags have to be within the radio range of the master tag so spreading them across huge areas is not possible. Instead of a tag being the master it is possible to use an anchor.

This makes it easier to have a computer attached, as the anchors are stationary, unlike the tags.

2.6 Unity introduction

As defined in [section 1.1](#), this project aims to create a location-based augmented reality game. This means the project has to have a game component - an application to display the objectives of the game, the play area, and the players. To create this, a game engine can be used, such as **Unity**.

A game engine is a piece of software that provides creators with the necessary set of features to develop games quickly and efficiently [21]. This means that a game engine is a collection of reusable components, abstracted away from the game developer. This can include tools to help with, for example, graphics, physics, networking or audio. These tools would expose certain functionality to a developer to make use of, and hide the specific implementation details for that functionality, ensuring the developer can focus on more pressing issues. Unity supports the C# language for development [18].

The Unity game engine supports development for different game platforms. Of particular interest to this project is the support for both **Android** and **iOS** devices, as well as **Google Cardboard** [19].

Unity was chosen for the development of the game aspect of this project since this facilitates that a greater amount of time can be spent on the other aspects of the project rather than the low-level details of game development, and it allows for easier inclusion of multiple platforms.

2.7 Networking

The following section will examine the different possibilities regarding transmitting player position data from the Pozyx tags to the Unity applications used to visualize the game.

2.7.1 Possible networking solutions

Unity will be used for the creation of the game aspect of this project as described in [section 2.6](#). Unity includes a proprietary networking solution known as UNet [20]. This solution allows developers to use a high-level API, giving access to commands that cover many common requirements for multiplayer games, without worrying about the low-level details.

Since the solution is developed alongside the actual game engine, it has a higher level

of integration with the Unity Editor and Engine, which allows for certain components and visual aids to aid the building of the game. As of the beginning of this project, the UNet solution has been deprecated for a while, and the Unity developers are actively working to create a new system to replace it.

The current UNet iteration is usable but will be removed in the future. Other third-party solutions for Unity-based games also exist, such as Photon Engine. Photon provides functionality for the developers to make use of to create multiplayer games in the same way as UNet, exposing higher-level functionality. Photon supports multiple platforms outside of just Unity, with both Android and iOS support [5].

The advantage of using a library that is built for Unity such as UNet or Photon is that it is easy to set up with the Unity engine compared to a solution built from scratch. UNet provides a higher level of abstraction with functions to control the networked state of the game, send and receive messages between server and clients and much more.

UNet and Photon also allow for "client-hosted" games that act like lobbies. So any client can host a game that other clients can connect to. This allows the clients to send and receive data between each other [20]. A disadvantage of these libraries is that they are generalized and thus would not be able to achieve the same efficiency as a custom solution tailored to the specific needs of the game.

ZeroMQ is also a possible solution. ZeroMQ is an asynchronous messaging library. It can carry messages across various transport formats and is available in many different programming languages [25]. It aims to be a high-performance library to be used in distributed or concurrent applications that are reliable. According to the getting started guide provided by ZeroMQ, certain issues tend to arise when developers attempt to create a networking solution using sockets [26]. These are:

- How to handle I/O?
- How are dynamic components handled? What happens if a component disappears temporarily?
- How are messages represented? Different sizes and different content can change representations
- How are messages that cannot be delivered immediately handled?
- Where should message queues be stored?
- How are lost messages handled?
- What if the network transport changes, for example, TCP to UDP?
- How do messages get routed? Can the same message be sent to multiple peers?
- How to write an API for another language?

- How to represent data such that it can be read between different architectures?
How much of this should be the messaging system's job?
- How do network errors get handled?

These issues are mostly applicable to general solutions that need to accommodate changing requirements or be reusable. However, for this project, not all of these issues are relevant. In terms of problems to overcome, this project should only be concerned with handling dynamic components, handling lost messages, routing messages and handling network errors.

If a player closes the game application it can lead to dynamic component issues. A message can be lost during the playing of the game. Messages should be delivered to all players to ensure that they all have the same information. Finally, a player might suddenly disconnect from the network.

The alternative to making use of a pre-existing solution is creating a custom solution. A custom solution entails a need to establish a familiarity with the required knowledge to construct such a solution. A custom solution would involve sockets, which are a network API that allows programs to communicate with each other [4].

2.7.2 Choosing a solution

There are certain pros and cons associated with both approaches of using either a pre-existing solution or a custom solution. Table 2.1 shows some of the considerations made when deciding an approach for this project. The criteria that were considered when evaluating which solution to use were:

- Customizability
 - How much we can customize the solution to our specific use case. If the customizability is low the project needs to be built around the networking solution, whereas if the customizability is high the networking solution can be customized to our needs.
- Requirements
 - How much knowledge about the subject is needed to use the solution.
- Optimization
 - How much optimization are we able to do ourselves if we use this solution.
- Learning reward
 - How much will we learn if we use this solution.

	Pre-existing Unity-based	Custom	ZeroMQ
Customizability	Consists of a set of pre-defined functionalities	Can have any functionality implemented	Has pre-defined functionalities, but these are lower level than a pre-existing solution
Requirements	Familiarity with the solution	Familiarity with the knowledge required to implement a usable solution	Needs familiarity with a mix of pre-existing and custom solutions
Optimization	Lower-level details are obscured, optimized for general use	Lower-level details are freely available, can be optimized for a specific purpose	Focuses on performance, but the solution is general
Learning reward	Most of it is already implemented, so the learning reward is low	Learning reward is high because we have to implement everything ourselves	Have to customize some of it ourselves, so we will have to get familiar with the subject

Table 2.1: A comparison of the pros and cons of the possible solutions

Based on these considerations, it was decided that a custom solution should be created to handle networking in this project. This choice was based on two major factors: the lack of transparency in a pre-existing solution as well as the need for fast communication and the opportunity to learn more about networking at a low level. For the game to be playable and enjoyable, the location data collected by the Pozyx system needs to be transmitted to all the clients as quickly as possible such that they always have an up to date view of the positions of the players.

To achieve this, it would be preferable to build a solution capable of performing the minimum amount of work as quickly as possible. Pre-existing solutions cannot be guaranteed to do the minimum amount of work as lower-level details are obscured from the developers. With a custom solution, the data sent across the network can be guaranteed to be exactly what is needed.

ZeroMQ was also a possible choice based on the performance needs, but its generalized approach concerning itself with reusability and issues unlikely to be a big factor in this project meant it was dismissed, in favor of a custom solution in which the problems defined in the previous section are handled. Additionally, we have not previously worked with networking at a low level, but wanted to learn more about

it, and decided that working without a framework would provide the best learning experience.

2.7.3 Introduction to sockets

In order to construct a network solution, a familiarity with the layers of a network is needed to gain an intuition of what sockets are. A common way to describe these layers is through the *open systems interconnection* (OSI) model for communication. This model is illustrated in Figure 2.8, along with approximate mappings of the technologies used for each layer.



Figure 2.8: An illustration of the OSI-layer, and the corresponding technology for each layer.

As shown, the relevant layers for the purpose of creating a networking solution through sockets is the third and fourth layers, transport and network. The network layer is handled by the IPv4 and IPv6 protocols, which will be discussed in subsection 2.7.6, and the transport layer is handled by either TCP or UDP, which is described in subsection 2.7.4.

The reason for the gap between TCP and UDP is to illustrate that it is possible to bypass this layer and use IPv4 or IPv6 directly [4]. Sockets provide the interface from the upper application layers to the transport layer. The upper layer handles details about the application, and the lower layers handle details relating to communication.

Programs that communicate across a computer network need an agreement on how those programs will communicate. This is known as a protocol. Generally, before defining the design details of the protocol, a decision should be made as to which program is expected to initiate communication. One way of defining this is through the client-server architecture illustrated in Figure 2.9. This split is used by most

network-aware applications [4]. The most common method of initiating communication when using the client-server architecture is to have the client initiate requests. This tends to simplify the protocol and the programs themselves [4].




Figure 2.9: An illustration of the client-server architecture with multiple clients

2.7.4 TCP and UDP

The following section introduces two different protocols for the transport layer - Transmission Control Protocol (TCP) and User Datagram Protocol (UDP). Both of these protocols use a network-layer protocol known as IP, which can be either the protocol IPv4 or IPv6.

TCP

TCP provides connections between clients and servers. A TCP client establishes a connection with a given server, then it receives or sends data to that server across the network and closes the connection. TCP provides reliability by making sure that every package that is sent is received by the receiver. When TCP sends data, it requires an acknowledgment from the receiver that the data has been received. If it does not receive such an acknowledgment, TCP automatically retransmits the data and then waits for an acknowledgment for the retransmission. After a certain number of retransmissions, TCP gives up.

Based on the implementation TCP will typically attempt to send data for 4-10 minutes. This does not guarantee that the receiver will receive the data, but the guarantee is that it will deliver the data if possible, or notify the user that the connection has been broken without an acknowledgment from the receiver. To know how long to wait for acknowledgments, TCP contains algorithms to estimate the *round-trip time* between the client and server dynamically. It also performs these estimations continuously, as the result can be affected by variations in the network traffic. TCP

sequences data by associating bytes and sequence numbers.

For example, if an application writes 2048 bytes to a TCP socket, it would be sent in two segments, with the first containing data with the sequence number 1-1024, and the second containing data with the sequence number 1025-2048. If they arrive in the wrong order, the receiving TCP reorders the segments based on the sequence numbers before passing the data to the receiving application. If the receiver receives duplicate data this can also be detected through the sequence numbers, and the duplicate data can be deleted.

TCP provides flow control by telling clients how many bytes of data can be accepted at any time, known as the window. The size of the window decreases as data is received, and increases as the receiver reads data from its buffer.

UDP

UDP is a simple transport-layer protocol [4]. An application writes a message to a UDP socket, which gets encapsulated in a UDP datagram, which further gets encapsulated in an IP datagram and then sent to the destination. A datagram is a self-contained entity of data carrying information to be routed from the source to the destination nodes without reliance on earlier exchanges between the nodes and the transporting network [6].

A UDP datagram is not guaranteed to reach its final destination, nor is it guaranteed that order will be preserved across the network, or that datagrams arrive only once. This means that the UDP protocol is unreliable. If a datagram is lost on the network and not delivered to the UDP socket, it will not be automatically retransmitted. UDP also does not provide an acknowledgment that datagrams were received, sequence numbers to ensure data can be ordered, *round-trip time* estimation or timeouts.

UDP has no notion of flow control, meaning a fast UDP sender can transmit data at a rate the receiver is unable to keep up with. As such, it does not provide the same reliability as TCP. If reliability is a requirement, it has to be built through features such as timeouts, retransmissions and adding acknowledgments from the receiving end. A UDP datagram has a length, which is passed to the receiving application along with the data. UDP is considered connectionless, as there does not need to be a long-term relationship between a UDP client and the server.

A UDP client can create a socket and send a datagram to a server, and then immediately send another datagram on the same socket to a different server. A UDP server can receive several datagrams on a single UDP socket, each from different clients.

Choosing between UDP or TCP

Windowed flow control might not be necessary for transactions where both ends agree on the maximum size of a request or a reply [4]. For this project, the most important type of message is the player location data. This message will always be

formatted in the same way, and as such, a maximum size can be agreed to mean the flow control aspect of TCP is not needed. Another aspect of TDP that is not needed for this project is the automatic retransmission of messages. While this can provide reliability, it is of no importance for the game.

The players of the game are only concerned with the most recent updates of their position. As such, if a message were to not be received, it would not make sense to continually delay subsequent messages to attempt to retransmit a message containing position data that is more and more likely to be outdated.

For the position data to be as recent as possible, the messages should be sent as frequently as possible. It is also not necessary to provide an acknowledgment that the message has been received. The receiving applications should simply update their locations to comply with the most recently received data. The sender should not be concerned that a message was received, it should just continue to send the next message, which is likely to be more recent. Duplicate messages also do not pose much of an issue. If the applications were to receive the same location data multiple times, it would not impact the overall functionality of the program, rather just the speed at which the next updates would be received.

UDP has no connection setup or teardown costs. UDP only requires two packets to exchange a request and a reply, whereas TDP requires about 10 packets [4], if a new TCP connection is established for each exchange. In terms of transaction time, the minimum time for a UDP request-reply is the round-trip time + server processing time, and the minimum time for TDP is $2 \times$ round-trip time + server processing time [4].

Because of the limited scope and uniformity of what is going to be transmitted via the networking solution, a lot of the features included in TCP are unnecessary. UDP is slightly faster because of its lack of reliability and other benefits but might require some extra work to implement some of the functionality that is missing when compared to TCP if this were to become necessary. Because of the reasons discussed, UDP seems to fit the needs of this project more than TCP and is the protocol chosen for the networking solution featured in this project.

An issue with the choice of UDP could present itself in that messages are not guaranteed to arrive in order. It could pose a problem if a player in the game received a message with recent location data, and then another message afterward with outdated data. This could cause the player objects in the game to be at positions in which they were in the past, but not currently in the present.

2.7.5 Introduction to UDP sockets

UDP is a connectionless, unreliable datagram protocol. Figure 2.10 shows an illustration of the client-server architecture using UDP. The client-side creates a socket and sends a request to the server as illustrated.

Once the request has been sent the client transitions to a state of awaiting a reply. Once the reply is received the client can send another request, or the socket can

be closed. The server side also creates a socket, and then binds the socket to a port. Once bound, the server can await a request from the client. When a request is received, the server processes it, and then sends it to the client after which it can return to awaiting requests.



Figure 2.10: An illustration of the client-server architecture with UDP

For the purpose of this project, it might not be optimal to have client request the server. The server does not need information from the client or acknowledgment, meaning the clients do not have to send messages. As such, it might be better to use a publisher-subscriber approach.

The server would then act as a publisher, constantly sending messages to the clients that would be subscribed to the publisher. An illustration of this concept using UDP can be seen in Figure 2.11.



Figure 2.11: An illustration of a publisher-subscriber architecture with UDP

Possible issues with using this architecture is that there would need to be a way to ensure all clients receive a sufficient amount of messages, and that the receivers are

not overloaded with messages.

Transmission in UDP

When working with UDP, there are three types of possible transmissions: Unicast, multicast or broadcast. Unicast is a one-to-one communication with a single source sending information to a single receiver, while a broadcast transmits the information to all nodes on a network.

The benefit of unicast is that it has been in use for a long time and utilizes well-established protocols, it is known from applications like HTTP, FTP, and Telnet [3]. For this project, however, the major drawback of unicast is that to transmit the message to multiple nodes, it has to send multiple unicasts messages addressed to each receiver, which also requires us to know the exact IP address of each destination device.

Looking at broadcast, we can instead transmit the information to all nodes on a network, which ensures that all nodes on the network receive the message. This could be useful if we only had the players of the game on our network, but since they might be connected to a larger network with a large number of clients, this could lead to the data being sent to clients that are not a part of the ongoing game session. Luckily, we have multicast in the middle of the two extremes, where you do not send from one client to another or one client to all others.

Instead, the data is sent to as many destinations as express an interest in receiving it [3]. This one-to-many approach seems suitable for this project since it would intuitively lead to better bandwidth utilization and does not require the receivers' addresses to be known.

2.7.6 IPv4 and IPv6

Both IPv4 and IPv6 provide packet delivery service for TCP and UDP. The major difference in IPv4 and IPv6 is the addresses they use. IPv4 uses 32-bit addresses, whereas IPv6, being a newer version, uses larger addresses of 128 bits [4]. IPv4 addresses are usually written as four decimal numbers separated by ".". This is known as *dotted-decimal notation*.

Each decimal number represents one of the four bytes of the 32-bit address. The first of the four numbers represent the address type. IPv6 addresses are usually written as eight 16-bit hexadecimal numbers [4]. The higher-order bits of the 128-address imply the type. For this project, this should not have a significant impact, and either can be used.

2.7.7 Format of data

As it is necessary to send data over the network, the format of the data must be decided upon. The network format for the field anchor, player position and goal position is `0xyyyxxxxiisstt`, where `t` is the type, `s` is the timestamp, `i` is the ID,

xxxx is the x position and yyyy is the y position. These formats can be seen on Table 2.2, Table 2.3 and Table 2.4. The data sent is almost identical, as the same information is required, the only difference between them is the type and the context of the IDs. The x and y position is the location of the player or anchor. For Table 2.4 the x and y position is the center of the goal, where the size of the goal is set to have a length equal to 20% of the shortest side of the playing field. The IDs for Table 2.2 are for each anchor. For Table 2.3 the ID 0 represents the ball and 1-4 represents a player and for Table 2.4 the ID identifies the team. The type is at the right hand side of the hexadecimal so that it is faster to decode and to switch on which type it is. After the type the timestamp is placed, which is used to detect outdated data. The timestamp is an integer value ranging between 0 and 255 such that when it hits 255 it resets to 0 to be able to continuously send timestamps, regardless of game length. The downside to this approach is since UDP does not guarantee delivery of messages, it may not go straight from 255 to 0 every time. To counter this, 0 to 20 is considered as newer timestamps than 235 to 255 so that the system is allowed to miss up to 20 messages near the end, to ensure that it does not end up being locked until the values reach the high range again.

It was decided to use a size of respectively 16 and 8 bits so that it is straightforward to work with in the C# language, because it is the same size as some variable types in C# and hereby easy to type cast. Given the highest number that can be in two bytes is $2^{16} - 1$, the maximum value for the x and y positions is 65.535.

POSY (y) (0-65.535)	POSX (x) (0-65.535)	ID (i) (0-3)	TIMESTAMP (s) (0-255)	TYPE (t) (0)
16 bits	16 bits	8 bits	8 bits	8 bits

Table 2.2: Format for field anchor position

POSY (y) (0-65.535)	POSX (x) (0-65.535)	ID (0-4)	TIMESTAMP (s) (0-255)	TYPE (t) (1)
16 bits	16 bits	8 bits	8 bits	8 bits

Table 2.3: Format for player positions.

POSY (y) (0-65.535)	POSX (x) (0-65.535)	TEAM ID (i) (0-1)	TIMESTAMP (s) (0-255)	TYPE (t) (2)
16 bits	16 bits	8 bits	8 bits	8 bits

Table 2.4: Format for the goal positions.

The goals scored has a different format of 0x1100sstt, where t is type, s is timestamp, 00 is the score of team 0 and 11 is the score of team 1. This format can be seen on Table 2.5.

TEAM1SCORE (1) (0-100)	TEAM0SCORE (0) (0-100)	TIMESTAMP (s) (0-255)	TYPE (t) (3)
8 bits	8 bits	8 bits	8 bits

Table 2.5: Format for goals scored.

2.8 Experiment with Pozyx

To determine the accuracy of the Pozyx tags an experiment was conducted. The primary goal of the experiment was to test the accuracy, but a secondary goal of the experiment was to determine the frequencies of updates for each tag.

2.8.1 Purpose

According to research regarding latency in VR [23] and interactive systems [14], the number of movement errors increase as the latency increases. An example of this is when a tag is moved, and the tag first sends a new position after 3 seconds, then the movement error could be high as the user does not know their exact position for these 3 seconds. According to experiments with VR, the motor performance and sense of body ownership start to decrease at latencies above 75 ms.

This means that the optimal results from these experiments would indicate that the Pozyx system can send position updates fast enough which allows the system to operate with a latency of less than 75 ms between their physical movement and the in-game reflection of this movement.

2.8.2 Setup

The tags were set to transmit with the highest bitrate with the longest preamble length and with the ranging mode set to precise [2].

According to the documentation, these values lower the tags update rate to about 9hz which is then divided by the number of tags used in the system as described in subsection 2.5.4. As the main goal of the experiment was to test the accuracy these settings were chosen to give the most accurate positioning. The settings can be changed at a later point to try to increase the update rate to a level where the latency for the users is acceptable.

The experiment was set up as shown on Figure 2.12. The experiment was conducted indoors on our campus in the building Novi 9. The anchors 0x632b and 0x676e were mounted on a wall 240 centimeters apart, and the remaining anchors 0x6738 and 0x676c were mounted on a bulletin board.

The number of centimeters accompanying the hexadecimal number of each anchor is the height at which the anchor was mounted during the experiment. Different heights were chosen as Pozyx documentation suggests that not all anchors should have the

same height [24]. The reason for this is the principle of geometric dilution of precision (GDOP), which can cause the error on range measurements to be amplified.




Figure 2.12: The setup of the experiment with the anchors and the height at which they were placed in the corners. The hexadecimal number is the anchor, and above that is the height at which the anchors were placed.

The fields were created by a whiteboard on which lines were drawn every 10 centimeters to know the actual position as seen on Figure 2.13. This whiteboard was moved intermittently to act as respectively fields 1, 2, 3, and 4.



Figure 2.13: The whiteboard with the drawn positions.

The procedure for the experiment was based on this whiteboard. The whiteboard would be placed in one of the fields defined in Figure 2.12, and the tags would be placed in certain positions on the board and record the accuracy with which the

position was reported. The tags would be placed in a position, remain there for five seconds, then be moved to the next position over the next five seconds and remain in this position for five seconds before being moved again. This procedure was repeated until a satisfactory number of measurements had been made. This amount was usually 5 measurements. Once one field had been tested, the whiteboard was moved to the next field, as seen on Figure 2.12, and the process was repeated.

2.8.3 Analyzing the data

In section 6.1 the data for the experiment can be seen. For each test the average grid was calculated, and the minimum and maximum values for x , y and z were found. The min and max values were noted, as they may be useful for future analysis. The z values were very inconsistent throughout the entire test, and due to this and the z -axis being largely irrelevant to the game, we will not focus on analyzing them. Because of this, it was decided to only calculate the deviance between the actual position and the measured position in the xy -plane.

Precision with 1 tag

When testing with only one tag the average of all deviations was 15.20 cm.

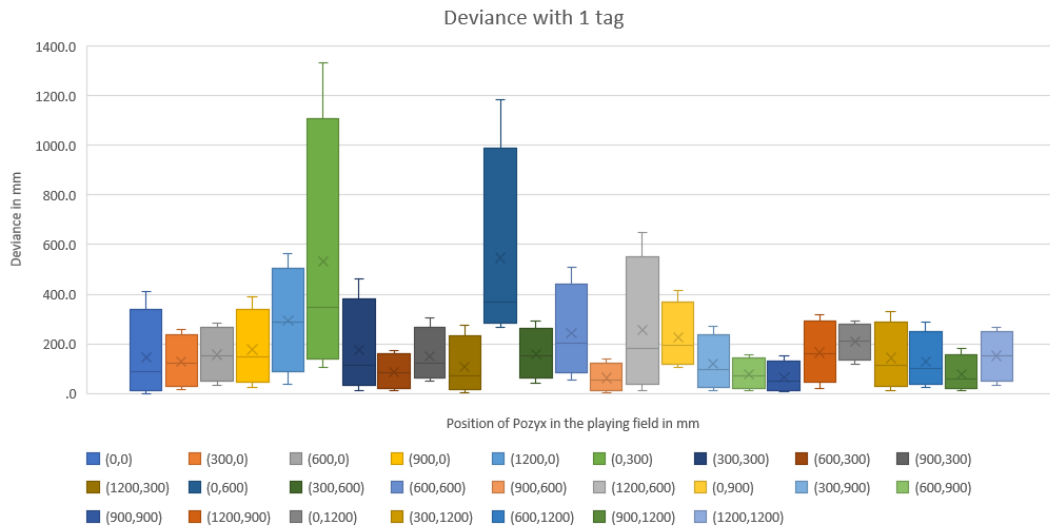


Figure 2.14: A box plot showing the dispersion of measured positions with 1 tag

As can be seen on Figure 2.14, some measurements can be quite far off the actual position. On the x -axis you can see the actual position of the tag in mm compared to the playing field. The y -axis shows the amount that the measurements were off the actual position in mm. Each measurement set has its own box and whiskers. The box shows the measurements within the first quartile and third quartile. Within the

box the X denotes the second quartile which is equal to the median and the flat line denotes the average value of all the measurements. Extending from the box is the lower and upper whiskers where the lower whisker shows the measurements below the first quartile and the upper whisker shows the measurements above the third quartile. The end of the lower whisker is the lowest deviance measured in the set, likewise the end of the upper whisker is the highest deviance measured.

Interestingly it is the measurements closest to the x -axis that are the most inaccurate, we will discuss why in [subsection 2.8.4](#). The measurement with the highest deviance was recorded when the tag was in position (0,3000) where the measured position was 1.3 meters off. The average deviance of all the measurements from that position was only 36.5 cm, however. The highest average deviance of all the positions was 39.5 cm in position (0,6000). When running with one tag the average amount of measurements were 12.368 per second.

Precision with 3 tag

For this part of the experiment, the accuracy was tested with three tags. The tags were positioned at three different x -coordinates. 1400, 1800 and 2200 were chosen as the x coordinates and the y coordinate was moved up by 300 mm each time starting from 0 and going up to 1200. The average amount of updates from all the tags per second were 1.83, less than 10 percent of the amount of updates when only one tag was used. The amount of updates were very different between the different tags, as the one with the highest update rate had an average of 3.4 per second and the lowest only 0.92.

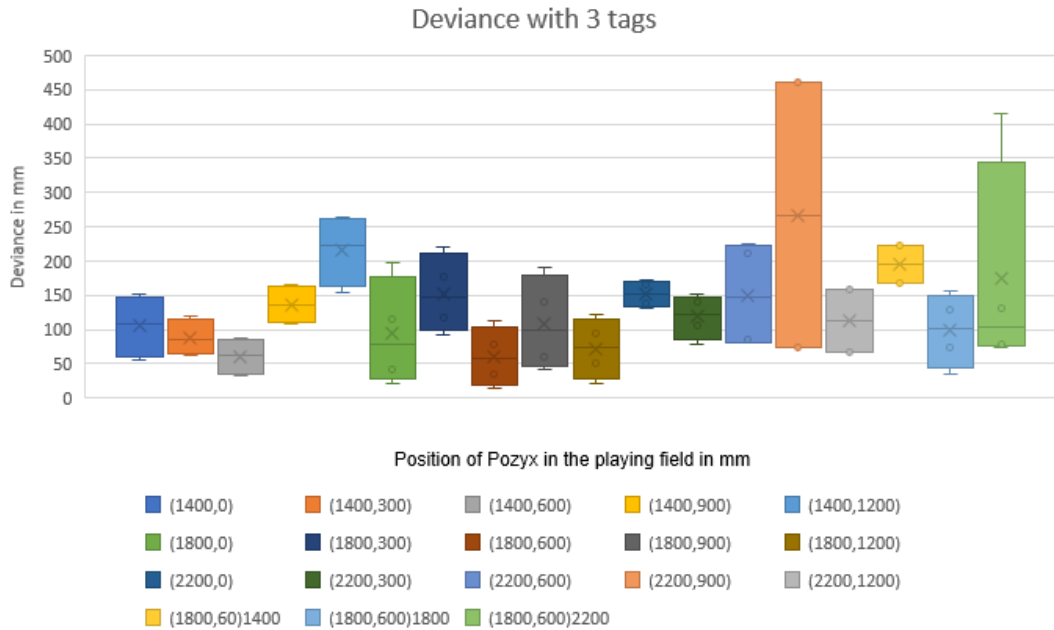


Figure 2.15: A box plot showing the dispersion of measured positions with 3 tags

On Figure 2.15 all the measurements from the 3 tag experiment can be seen. The 3 last sets of measurements are from when all the tags were placed close to each other at the same point. The trailing number after the coordinate is which tag the set of measurements belong to, (1800,600)1400 denotes the tag used for the measurements along the line where x equals 1400.

This was done to see if they would affect the precision of each other. Based on the data it seems they might have a small effect on each other but not enough for it to be a problem. The biggest deviance in the test was only 46 cm off compared to 1.3 meters in the 1 tag test which seemed interesting, but could be because the measurements in the 3 tag experiment were made in the middle of the playing field compared to at the edge in the 1 tag experiment.

Tag 26467 This tag was positioned 1400 mm along the x -axis. The overall average deviation for this tag was 11.9 cm and it made 1.16 measurements per second. No measurements were further away than 26.3 cm from the actual position and the biggest average deviation for a set of measurements was only 22.26 cm.

Tag 26895 This tag was positioned 1800 mm along the x -axis. This tag was only 9.2 cm off on average and had by far the highest average update rate at 3.4 Hz. Like the previous tag it also did not have any measurements with large deviances. The furthest measurement was 22.1 cm off and the worst average deviation for a set was 15.09 cm. This was by far the best performing tag in this experiment regarding both

precision and update rate.

Tag 24622 This tag was positioned 2200 mm along the x -axis. The overall average deviation for this tag was 14.48 cm and an update rate at 0.92 hz. For the first three points the results seemed fairly accurate, however, the grid (2200, 900) had a large deviation of 22.43 cm. That set also had the measurement with highest deviance with 46.0 cm.

Precision with 5 tag

The tags were positioned at five different x -coordinates. 1400, 1600, 1800, 2000 and 2200 were chosen as the x coordinates and the y coordinate was moved up by 200 mm each time starting from 0 to 1200 mm.




Figure 2.16: A box plot showing the dispersion of measured positions with 5 tags

Like the experiment with 3 tags, seen on Figure 2.15, there were no huge outliers like in the first experiment where some measurements were off by more than a meter. But two instances were still seen as further off than the average.

Tag 24622 This tag was positioned 1800 mm along the x -axis. The average deviation for this tag was 12.99 cm and the average update rate was 0.83 Hz. This was the tag that had the worst measurement on coordinates (1800, 0) which was off by

65.1 cm and the worst average deviance of a measurement set with 31.7 cm. Besides that the other sets were among the most accurate.

Tag 26467 This tag was positioned 1400 mm along the x -axis. For this tag, the average deviation was 12.75 cm and the average amount of data points per second was 0.6. At 1400, 200 the tag only returned one measurement and at (1400, 600) there were only two measurements completed during the 5 second period the tag was in that position. The tag was overall very precise with all measurements within 23.6 cm of the actual position.

Tag 26895 This tag was positioned 1600 mm along the x -axis. Tag 26895 had an average deviation of 12.39 cm and an update rate at 2.11 Hz. This update rate was much higher than tag 24622, 26467 and 27001, and equal to 26901 even though they all had the same settings.

Tag 26901 This tag was positioned 2200 mm along the x -axis. The average deviation of tag 26901 was 14.74 cm, and it had an update rate of 2.11 Hz aswell. This tag had a few spikes when measuring on coordinate (2200, 200), where the average deviance was 20.99 cm and the worst measurement was 54.6 cm off.

Tag 27001 This tag was positioned 2000 mm along the x -axis. The average deviation of this tag was 10.46 cm and the average update rate was 1.23 Hz. This tag was pretty average with no big deviations. Its worst measurement set had an average deviation of 18.34 cm.

2.8.4 Possible influences on the test

One thing that could have affected the tags is that a whiteboard was used as a measure for positions. As there is metal in the whiteboard this could have affected the precision of the tags. Metals are conductors, which can lead to the signal having less power and reduced range, and the signal might spend extra time trying to get through the metal. Since Pozyx positioning relies on calculating the time of flight, having the signal spend extra time traveling reduces accuracy [22].

If we look at the z coordinate in the test experiment data, it can be seen that z fluctuates a lot more than the x and y coordinates. The max z coordinate for multiple experiments is often more than 3 meters and the average z coordinate is also often 1 meter higher than the actual coordinate. This could be due to the height difference between the anchors not being sufficiently large which can affect Pozyx' ability to measure the height of the tags.

During the test with 1 tag, the 0 value of the x coordinate coincided with the wall. This could have been an influence on the positioning result in that it could increase

uncertainty, as it was difficult to center the tag over the x coordinate since the coordinate collided with the wall. Also, the wall could have created interference with the signal and there might have been metal or wires going through it which could have affected the signal.

While the experiment was running the tags could throw errors instead of giving a position if something went wrong. Mainly during the multi-tag tests, some tags would report back with errors such as unable to get firmware version, flash memory corrupted or an error message saying there was no error. At that point in the experiment, it was not clear whether it was a hardware error with some of the tags or if it was code related. The tags would throw these errors randomly and then measure their position normally at the next update. This resulted in some of the tags having periods of update rates of less than 0.2 Hz. It was decided that further investigation into this issue should be conducted at a later point.

2.8.5 Conclusion on the experiment

Our conclusion of this experiment is that the precision is satisfactory, but the update rate is not optimal as that does not meet the standards of the refresh rate needed with the current settings.

We also need to consider the spikes in the coordinates, such that the user does not jump around on the screen, even though the user might not be moving. A good solution would be to use an algorithm to correct the data.

The results of this experiment shows that a three dimensional game is not viable with the errors that we had on the z coordinate. This is most likely due to the height difference between the anchors only being 20 cm. If we, later on, choose to make use of the z coordinate, then it is necessary to conduct another experiment with the Pozyx with different, larger height differences, but as it currently is only a two-dimensional game it is not necessary.

It also seemed like some tags were not able to send as many measurements as others. Consistently throughout the experiments tag 26895 was the best tag regarding update rate, sometimes having an update rate almost 3 times faster than the other tags. This was caused in part due to some tags seemingly having a slower update rate, but also because some of them would report errors instead of measurements randomly. This is something that will be investigated.

2.9 Sprint 1 conclusion

This section concludes the preceding chapter on sprint 1. It will conclude what knowledge was gained based on the sections of the chapter, and discuss a retrospective

of the sprint as a whole, and what changes were made to the progress for upcoming sprints.

2.9.1 Overview of completed tasks

section 3.1 outlined the goals of the sprint. The sprint focused on exploring the project idea. Wireframes were constructed in section 2.3 in order to create a more uniform vision of the project. These generated a vision for how the users would interact with the game, and how the game aspect would look.

In section 2.4 a rudimentary overall architecture for the project was outlined. We defined an architecture in which data would be transferred from the Pozyx components to a host and then to the players of the game. section 2.5 gave an introduction to the Pozyx technology to be used for the tracking aspect of the project. Several alternatives for finding anchor locations were proposed, and UWB was defined as the best solution. section 2.6 gave an introduction to the game engine to be used for building the project. section 2.7 explored the possibilities for communication between the different parts of the architecture, settling on using a custom UDP based solution. Finally, section 2.8 defines an experiment conducted in order to determine the accuracy of the Pozyx system, concluding that the accuracy was acceptable for the use of this project.

2.9.2 Retrospective on the process

Initially, the project made use of a more scrum-inspired process, which consisted of smaller sub-sprints of one week composing the larger sprints defined as sections in this report. Each week had a backlog of tasks to be completed that week, evaluated on a story point-based system. During this sprint, this was changed to the process outlined in 2.2. The shorter sprints were phased out, and instead, a more constant backlog was kept, from which to choose tasks. We held a retrospective on this process at the end of the sprint. The developers liked not having the smaller sprints and decided to do reviews in pairs going forward. To ensure pull requests would be reviewed quickly, a rule to check for pull requests every morning was set in place. The waypoints were assigned to tasks in the process was questioned, as the way it was done lead to some bias towards tasks that took a short amount of time to complete. It was decided that creating a weight to influence the tasks in a way to avoid this bias was a good idea.

Chapter 3

Sprint 2

3.1 Sprint goal and introduction

The goal for this sprint is to explore the networking side of the project and implement initial versions of the UDP client and server based on what was learned in [section 2.7](#). Additionally, the first version of the game should be created, which should allow players to see the game through a VR headset and show the players moving based on the positions of the tags.

3.2 Deployment Diagram

In [section 2.4](#) the different components of the system were introduced in [Figure 2.6](#). A deployment diagram is constructed in this section in order to further elaborate on the different components.



Figure 3.1: A deployment diagram for the system.

A deployed system will contain three nodes:

- Pozyx system
- Host
- Mobile Clients

Nodes are represented by cubes, and are entities that execute components. The Pozyx system node contains two components. These are the anchors and the tags needed to generate positional data with the Pozyx hardware. The system contains multiples of anchors and tags. These Pozyx components generate artifacts in terms of positional data for the location of the players and the ball. Anchors and tags are associated through a UWB connection in order to generate these artifacts. The host node is dependent on the positional data artifact, as the component receives the positional data and transforms it for communication. Positional data is transferred from the Pozyx system to the host through a USB association. One anchor is associated with the host application, and this anchor is responsible for collecting data from all tags. The host application component contains a user interface component, which is the interface that the person using the host component will interact with, in the form of a console application for the early version of the system. The host node is

associated with the clients through a UDP connection in order to communicate both the positional data and the configuration data to perform game setup. As such, the mobile client application component is dependent on the host application component, since it requires the configuration data. The mobile clients also contain user interface components, which is how the users interact with that part of the system. This user interface is the virtual playing field generated in Unity that the users view in the headset.

3.3 Accessing games on the network

Different methods of transmission were discussed in [subsubsection 2.7.5](#), and multicasting was selected as the optimal solution for this project. In order to make proper use of this in an eventual deployment of the game, it would not make sense for the users to input the IP address to which they want to connect. As such, a way to access hosted games without this needs to be implemented. In order to support multiple games being played at the same time, the system also needs to properly make use of multicast groups based on the game being joined. Clients should only receive data from one specific host in their specific multicast group.

To do this, the game must include some form of LAN discoverability. One way of achieving this is to have clients searching for a game broadcast a message to available hosts via the LAN. The hosts then reply, if they are available, with data for a multicast group, and the client can then join that group. Another method could be to make use of Internet Group Management Protocol (IGMP), which is a communications protocol used on IPv4 networks to establish multicast groups. The IP address 224.0.0.1 is a notable IPv4 address reserved for IP multicasting, which is the multicast group address for all hosts on the same network [13]. All hosts should join that group on start-up, and clients could then message the group that they are looking for an available game.

Another way would be to have the host continually broadcast messages that it is available for players. This would, however, lead to the host needing to repeatedly send messages as long as the game it is hosting is not filled with players or has not started. This would likely be a less elegant solution than having players searching for hosts message a couple of times.

Having the players send a message to the multicast group address for all hosts is the preferable solution, as it will avoid unnecessary overhead caused by broadcast messages being sent to unrelated machines on the network. In order to select a game to join as a client, the game would need a game selection screen. [Figure 2.3](#) and [Figure 2.4](#) show the initial prototypes for joining a host based on their IP and waiting for the game to begin. If a lobby were to be implemented, [Figure 2.3](#) would need to be updated, and there would need to be an extra step before [Figure 2.4](#) could be

shown. Figure 3.2 and Figure 3.3 show new prototypes for this purpose. Choosing a lobby on Figure 3.3 would lead to Figure 2.4.

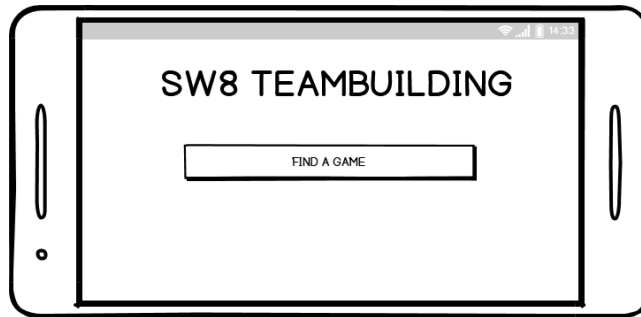


Figure 3.2: Prototype of the game menu when players can choose a game to join.




Figure 3.3: Prototype of a lobby menu. Available games are shown as well as the number of players. Players can choose one to join.

3.4 UDP implementation

In this section, we will delve into the implementation details of the UDP client and host. As described in section 3.2, the system will consist of a host computer which receives positional data from Pozyx and transmits it to the players using IP multicasting.

Host

The main functionality of the host is to transmit positional data from the Pozyx tags to the clients via multicasting. The first setup step of this is to connect the host to an IP address on the network dedicated to multicasting, as described in section 3.3. For the time being, this is hardcoded to be 224.3.29.71:10000. After this initial setup, the host will continuously run through a loop where it updates the ball position, player positions and increments the timestamp. Each update consists of three steps: Getting the positional data from Pozyx, transforming it to fit the format described in subsection 2.7.7 and transmitting the data to the clients, as seen on Listing 3.1.

```

1 def update_ball_position(self):
2     """Broadcasts the updated ball position"""
3     ball_tag = self.setup.ball_tag
4     position =
5         self.multi_tag_positioning.get_position(ball_tag)
6     message =
7         self.formatter.format_player_position(self.time_stamp,
8         0, position.x, position.y)
9     self.multicast_sender.send(message)

```

Listing 3.1: Updating ball position

Client

Like the host, the client joins the multicast group on the specified IP. While the host continuously sends positional data, the client's job is to receive and react to the data that is being sent to the multicast group. This is done by creating an asynchronous callback, which ensures that the data receiver is called whenever data is being transmitted, which is seen on the first three lines of Listing 3.2. Once the **Receive** function is called, it saves the received bytes to an array and starts listening for new data before it starts working with the data. This was done to ensure that new messages would not be blocked if it takes too long to handle the data. In the current version, the receiver will simply log the data to the Unity console, but in future versions, it will be able to call the appropriate functions based on the type of data it has received.

```

1 public void StartListening()
2 {
3     uClient.BeginReceive(new AsyncCallback(Receive), null);
4 }
5
6 public void Receive(IAsyncResult res)
7 {
8     // Represents a network endpoint as IP address and port
9     // number.
10    IPEndPoint RemoteIpEndPoint = new
11        IPEndPoint(IPAddress.Any, portNumber);
12
13    // Receives the message as an array of bytes, then ends
14    // communication with the remote endpoint.
15    Byte[] receiveBytes = uClient.EndReceive(res, ref
16        RemoteIpEndPoint);
17
18    // Restarts communication again to receive a new datagram.
19    uClient.BeginReceive(new AsyncCallback(Receive), null);
20 }

```

```
17     // The bytes that were received are converted to a string,  
18     // which is written to the unity debug log.  
19     string returnData =  
20         System.Text.Encoding.ASCII.GetString(receiveBytes);  
21     datagramMessage = returnData;  
22     datagramSender = "Address: " +  
23         RemoteEndPoint.Address.ToString() + ", port: " +  
24         RemoteEndPoint.Port.ToString();  
25 }
```

Listing 3.2: Receiving data from host

3.5 Current product

To give an overview of the progression of the project this section will provide an overview of what has been created during the sprint. This overview will be split into two categories, to reflect the structure of the project: Networking and game.

3.5.1 Networking

The network aspect has been a big focus this sprint and has led to a better understanding of how the data should flow through the system (section 3.2), and how the knowledge from sprint 1 about networking can be used to automatically find on-going games on the local network (section 3.3).

From an implementation perspective, this sprint has introduced an initial version of the UDP host which is capable of transmitting Pozyx location data from the host computer to the game clients. For the host, an initial console-based setup has been created, where the user can input the number of players, the location of the anchors as well as specify which tags are used in the game. Additionally, the host can automatically re-order the list of entered anchors to ensure that they are sent to the clients in a clockwise manner, such that Unity can create a playing field mesh based on the coordinates.

3.5.2 Game

For the game aspect of the project, the most crucial parts have been implemented this sprint:

First of all, the game has been set up to support VR glasses by splitting the game view into two halves, one for each eye, as seen on Figure 3.4.




Figure 3.4: The first implementation of VR

The game supports moving the players upon receiving new information from the host, but the connection between the two was not coupled together in this sprint. Finally, an algorithm was implemented to generate the goal zones and ensure that they are placed fairly on the playing field.

3.6 Retrospective on the process

The process required some unexpected pivots in the second sprint due to the coronavirus, which led to the university being locked down for a period.

The daily group work was moved from the physical group room to a virtual setting facilitated by Discord. One of the major differences was that it was previously possible to sit down and do the pair reviews in the morning, which meant that reviews could be completed quickly and easily. While we tried to still get the pair reviews done first thing in the morning, there was a noticeable delay before the new features were reviewed and merged. Another minor change was how we assigned points (cost, reward, priority) to tasks. This worked out fine online, as people could simultaneously post their chosen number in a chat, as opposed to showing a number of fingers while sitting around a table.

Since most of the project work is managed on Jira and completed by individual group members, there was not a noticeable change in the amount of work getting done. A side effect of working with only voice chat was occasional difficulty in reaching the other members when their help was needed since they might be temporarily away from the computer, or too distracted or focused to be paying attention to the voice chat. This also resulted in some information not getting spread to the entire group, as not everyone may have been listening while a discussion was happening, which led to minor confusion in further discussions on the subject.

Like at the end of sprint 1 (section 2.9), a retrospective was conducted by the anchor to reflect upon how the process was working out. In addition to discussing the progress of the project, the following points were brought up:

How does the process with Jira work?

Currently, it seems like the challenger is the only person adding suggestions to the backlog, after some discussion it turned out that multiple members thought that this was intentional, and did not know that they were also supposed to contribute with their suggestions for the project. This has been clarified, and all members are now aware that they are fully encouraged to add suggestions to the Jira when they come up with ideas for the project.

Should we use pair programming more?

Right now, the only work done in pairs is the pair reviews. It was decided that utilizing pair programming for larger programming tasks would be beneficial to decrease the amount of time spent on it and allow for more perspectives on implementation. Whether or not a task should be done with pair programming will be decided as a part of the task discussions that take place after the daily standups if there are new tasks in the suggested column.

How do we give better estimates about when a task is done in daily standups?

Generally, the biggest reason that it is difficult to estimate how much remains of a given task, is that the definitions of done are not precise enough. To combat this, it was decided to remove the prioritization of tasks where we would assign a reward, cost, and priority. Instead, the discussion will be about what a good and specific DoD is for the given task to ensure that everyone is on the same page. The hope with this approach is that discussing the DoD will give new perspectives to a given task, and shift the focus from how long it will take to complete the task to exactly what needs to be done. A possible by-product of this shift of focus is that the in-depth discussion will result in discovering new tasks that need to be completed.

How do we feel about daily standups?

Everyone is generally pleased with the daily standups, but there is a tendency to focus more on what has already been done rather than what is currently being worked on. To prevent this, a new format for the meetings has been proposed:

- What have I been doing? (short)
- What am I working on now?
- When will my task be completed?

- What is my current challenge?
- Do I need help or reviewers?

After each member has presented these four points, we will go through the new tasks in the suggested column and define a DoD for them. Finally, we went through all tasks that had been completed in the sprint to ensure that everything that had been implemented was also documented in the report, to ensure that the knowledge is shared between all members of the group.

Chapter 4

Sprint 3

4.1 Sprint 3 goals

The main goal for sprint 3 in terms of implementation is to achieve a minimum viable product(MVP). To do this, the game should be able to receive positional data from the Pozyx system. This data should then be formatted and used to update the positions of the players. On top of this, the host should be able to calculate when a player has entered the opposition's goal with the ball, and notify the players that a goal was scored. Finally, a win condition should be implemented such that a game can end, such as one team reaching three goals scored.

4.2 Dead reckoning

The accuracy of the representation of the players as well as the ball is important for the game to function well. Since these positions are sent over the network, they might not always be the most recent positions, since there could be delays or packages could be lost.

Dead reckoning is a technique that is used to predict where a game object is at a given time based on its last known position, velocity and acceleration. It is often used for networking games where packages about the object's kinematic state are continuously being sent from the server to the clients. The kinematic state of an object includes its position, velocity, acceleration, orientation and angular velocity [15]. If the client misses a package or they are not being sent often enough to have a new update for each frame, the representation of the object will jerk across the screen instead of having smooth and consistent movement. That is where dead reckoning can be used to predict an object's movement to make it appear more believable. For the game we are building, it is important that each player's and the ball's positions are very accurate for the game to function properly.

It is not possible to have a new update for each frame, zero packet losses or zero

latency. Therefore dead reckoning is needed to achieve a believable representation of the ball and the players' movement [15].




Figure 4.1: Linear player movement

Figure 4.1 shows an example of a player's movement in the game. In this example, dead reckoning would be a linear problem where the position, velocity and acceleration can be used to predict where the player will move to in the future. The dead reckoned position for a specific time Q_t in this example can be calculated by:

$$Q_t = P'_0 + V'_0 T + \frac{1}{2} A'_0 T^2$$

where P'_0 is the position, V'_0 is the velocity, A'_0 is the acceleration and T is the time [15].




Figure 4.2: A new update about the player's state is received from the server

In Figure 4.2 we receive a new update about the player's kinematic state. This results in conflicting realities when compared to the dead reckoned position calculated from the previous example. In the new update the player has turned right and a new dead reckoned position must be found, but this time it becomes a lot trickier since a believable curve must be created from where we thought the player would be and where we estimate the player will be based on the new information. Representing the new position P'_0 immediately would result in the player warping across the screen which would not be ideal. Instead, the player is represented at P_0 where the player was thought to be, and the player should then move towards a new dead reckoned position, calculated with the information from the new update [15].

4.2.1 Projective Velocity Blending

To calculate the curve that the player's movement needs to follow requires a good algorithm that is not too CPU intensive. The algorithm must work well for a segment of a curve that is passing through two points, those being the player's current position P_0 and the estimated future location P'_1 . The recommended approach for this problem is projective velocity blending [15].

Projections are created for the current kinematic state and the last known kinematic state and then these are blended together.

$$P_t = P_0 + V_0 T_t + \frac{1}{2} A'_0 T_t^2 \quad (\text{Projecting from last dead reckoned position})$$

$$P'_t = P'_0 + V'_0 T_t + \frac{1}{2} A'_0 T_t^2 \quad (\text{Projecting from last known})$$

$$Q_t = P_t + (P'_t - P_t) \hat{T} \quad (\text{Combination of the two})$$

Here, T_t is how much time has elapsed since T_0 which is the time at the last known position, which would be the beginning of the curve. T_1 is the time when the player reaches the end of the curve. \hat{T} is a normalized time value where $0.0 \leq \hat{T} \leq 1.0$ and should represent how far along the curve the player is.

From these calculations we get the dead reckoned location Q_t at the specified time. The reason that A'_0 is used in both projections is that it will converge to the player's true path much faster and it reduces oscillation, compared to when using A_0 when calculating P_t . However, these calculations will still give inadequate results since the player's movement will have bad oscillations. These are caused by the changes in velocity, V_0 and V'_0 , when new updates are received from the server. To account for this, a linear interpolation between the old velocity and the last known velocity is computed, which creates a blended velocity V_b . This velocity is used in the projection from where the player was, when a new update is received. This is what is known as *projective velocity blending* [15].

$$V_b = V_0 + (V'_0 - V_0) \hat{T} \quad (\text{Velocity blending})$$

$$P_t = P_0 + V_b T_t + \frac{1}{2} A'_0 T_t^2 \quad (\text{Projecting from last dead reckoned position})$$

$$P'_t = P'_0 + V'_0 T_t + \frac{1}{2} A'_0 T_t^2 \quad (\text{Projecting from last known position})$$

$$Q_t = P_t + (P'_t - P_t) \hat{T} \quad (\text{Combination of the two})$$

This should reduce the oscillations in the player's movement significantly.




Figure 4.3: Shows the the curve for the dead reckoned position Q_t when making use of projective velocity blending

In Figure 4.3 we see the red curve that the player should follow to reach the dead reckoned position Q_t . P_0 is where the player is currently represented. P'_0 is the recent position received from the server. The green curve is the one that the player actually follows.

This technique would be ideal to implement in the game if we find that the representation of the players' movement is inconsistent and in need of improvement.

Chapter 5

Sprint 4

Chapter 6

Appendix

6.1 Results for experiment

6.1.1 Experiment with one tag

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(0,0)	(8.8, 0.2, 113.5)	(0.0)	(34.0)	(-23.1)	(11.0)	(0.0)	(346.5)
(30,0)	(37.0, 9.1, 180.8)	(26.9)	(51.2)	(0.6)	(21.3)	(74.7)	(336.3)
(60,0)	(70.8, 9.9, 221.1)	(55.7)	(84.3)	(-3.6)	(28.1)	(84.0)	(325.4)
(90,0)	(100.2, 5.7, 263.1)	(83.4)	(108.9)	(-37.3)	(21.5)	(79.3)	(322.6)
(120,0)	(135.6, -22.6, 270.3)	(116.1)	(172.1)	(-56.3)	(14.0)	(48.2)	(305.0)
(0,30)	(24.2, 36.1, 199.0)	(11.6)	(57.0)	(25.6)	(67.7)	(28.4)	(334.1)
(30,30)	(66.9, 30.7, 193.5)	(56.8)	(77.1)	(19.3)	(40.2)	(66.0)	(324.3)
(60,30)	(100.6, 24.6, 292.2)	(83.1)	(109.5)	(0.2)	(33.0)	(88.2)	(318.8)
(90,30)	(121.7, 35.5, 291.4)	(104.8)	(136.8)	(14.1)	(53.3)	(64.3)	(313.3)
(120,30)	(31.2, 22.9, 242.8)	(9.0)	(107.1)	(-49.2)	(57.0)	(24.5)	(357.6)
(0,60)	(29.6, 39.1, 128.7)	(-11.9)	(105.7)	(6.3)	(113.2)	(-23.1)	(356.6)
(30,60)	(39.2, 60.1, 116.8)	(5.7)	(53.0)	(49.7)	(83.8)	(69.9)	(334.9)
(60,60)	(64.9, 77.5, 255.2)	(47.0)	(89.0)	(52.4)	(101.9)	(45.9)	(339.3)
(90,60)	(91.2, 64.0, 120.4)	(80.7)	(98.8)	(56.8)	(70.7)	(70.6)	(320.3)
(120,60)	(110.3, 52.3, 127.4)	(74.3)	(143.9)	(13.8)	(76.8)	(62.2)	(314.4)
(0,90)	(17.0, 90.3, 122.4)	(4.9)	(30.1)	(73.8)	(122.9)	(47.9)	(334.1)
(30,90)	(37.7, 90.2, 242.5)	(19.4)	(54.0)	(69.7)	(102.2)	(61.4)	(330.6)
(60,90)	(65.6, 89.7, 105.4)	(54.1)	(75.3)	(83.4)	(98.8)	(82.7)	(317.6)
(90,90)	(88.0, 91.8, 92.2)	(80.9)	(95.7)	(75.2)	(99.2)	(82.1)	(304.4)
(120,90)	(107.5, 84.8, 90.8)	(88.1)	(126.9)	(73.2)	(95.0)	(65.0)	(314.8)
(0,120)	(20.8, 119.4, 67.5)	(11.6)	(28.9)	(112.5)	(131.4)	(59.5)	(74.4)
(30,120)	(33.7, 114.9, 213.4)	(15.9)	(53.2)	(87.8)	(138.6)	(52.7)	(333.1)
(60,120)	(66.0, 127.4, 281.4)	(45.1)	(83.2)	(111.4)	(143.1)	(82.2)	(310.9)
(90,120)	(89.4, 120.7, 92.8)	(76.3)	(98.3)	(111.9)	(132.0)	(76.5)	(314.2)
(120,120)	(134.1, 121.8, 86.3)	(112.2)	(145.4)	(100.2)	(133.4)	(61.9)	(337.4)

Table 6.1: Table with grids for experiment with 1 tag. Tag: 26895

Position	Amount of data points	Average deviation
(0, 0)	58	9.92 cm
(30, 0)	65	12.29 cm
(60, 0)	54	15.76 cm
(90, 0)	65	14.74 cm
(120, 0)	62	28.85 cm
(0, 30)	59	36.56 cm
(30, 30)	64	11.65 cm
(60, 30)	61	8.40 cm
(90, 30)	67	12.82 cm
(120, 30)	53	7.95 cm
(0, 60)	65	39.56 cm
(30, 60)	61	15.32 cm
(60, 60)	60	20.34 cm
(90, 60)	65	5.39 cm
(120, 60)	66	18.90 cm
(0, 90)	65	20.10 cm
(30, 90)	63	10.28 cm
(60, 90)	59	7.23 cm
(90, 90)	59	5.39 cm
(120, 90)	65	15.23 cm
(0, 120)	65	21.12 cm
(30, 120)	57	11.71 cm
(60, 120)	59	11.10 cm
(90, 120)	67	5.87 cm
(120, 120)	62	15.26 cm

Table 6.2: Average deviation for each position with tag 26895

6.1.2 Experiment with three tags

Tag 26467

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(140,0)	(141.7, -8.4, 268.2)	(134.5)	(147.7)	(-15.1)	(1.6)	(79.1)	(325.5)
(140,30)	(140.1, 31.4, 124.0)	(134.1)	(150.4)	(23.1)	(39.1)	(77.5)	(306.6)
(140,60)	(133.7, 59.1, 162.7)	(131.3)	(137.0)	(57.3)	(61.2)	(87.5)	(310.0)
(140,90)	(126.5, 88.5, 90.8)	(123.8)	(129.2)	(86.7)	(91.9)	(87.8)	(96.1)
(140,120)	(118.6, 117.6, 86.1)	(113.7)	(128.9)	(109.3)	(121.0)	(80.2)	(91.7)

Table 6.3: Table with grids for experiment with three tags. Tag: 26467

Position	Amount of data points updates	Average deviation
(140, 0)	7	10.8 cm
(140, 30)	6	8.79 cm
(140, 60)	6	6.47 cm
(140, 90)	5	13.72 cm
(140, 120)	5	22.26 cm

Table 6.4: Average deviation for tag 26467**Tag 26895**

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(180,0)	(184.0, -2.0, 239.3)	(172.6)	(192.7)	(-10.7)	(18.2)	(86.6)	(304.9)
(180,30)	(184.9, 18.4, 165.5)	(174.5)	(202.1)	(9.3)	(29.1)	(79.7)	(298.1)
(180,60)	(179.9, 57.8, 155.3)	(173.7)	(187.1)	(49.6)	(67.5)	(84.9)	(307.2)
(180,90)	(174.6, 95.1, 198.1)	(167.0)	(187.0)	(85.8)	(104.6)	(85.7)	(320.5)
(180,120)	(175.4, 116.8, 82.6)	(168.6)	(180.5)	(108.5)	(122.7)	(74.0)	(94.5)

Table 6.5: Table with grids for experiment with three tags. Tag: 26895

Position	Amount of data points updates	Average deviation
(180, 0)	17	8.46 cm
(180, 30)	17	15.0 cm
(180, 60)	17	5.45 cm
(180, 90)	17	9.91 cm
(180, 120)	17	7.09 cm

Table 6.6: Average deviation for tag 26895**Tag 24622**

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(220,0)	(234.1, 5.4, 324.5)	(232.7)	(236.1)	(2.8)	(8.2)	(323.9)	(325.4)
(220,30)	(229.5, 28.2, 264.3)	(216.7)	(234.4)	(22.3)	(37.1)	(90.9)	(298.2)
(220,60)	(223.6, 46.1, 214.6)	(215.5)	(231.3)	(40.5)	(53.4)	(105.9)	(288.1)
(220,90)	(203.4, 78.6, 237.7)	(190.0)	(213.4)	(55.1)	(94.2)	(67.2)	(329.2)
(220,120)	(216.4, 119.1, 162.7)	(204.2)	(226.4)	(113.5)	(124.5)	(80.8)	(322.4)

Table 6.7: Table with grids for experiment with three tags. Tag: 24622

Position	Amount of data points updates	Average deviation
(220, 0)	5	15.18 cm
(220, 30)	7	11.76 cm
(220, 60)	5	15.30 cm
(220, 90)	3	22.43 cm
(220, 120)	3	10.08 cm

Table 6.8: Average deviation for tag 24622

6.2 Experiment with five tags

Tag 24622

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(180,0)	(173.8, -30.0, 125.8)	(168.3)	(180.1)	(-65.1)	(0.0)	(93.7)	(277.2)
(180,20)	(182.1, 25.4, 160.4)	(175.2)	(189.9)	(16.9)	(32.3)	(79.4)	(309.8)
(180,40)	(183.5, 41.3, 90.9)	(179.6)	(185.5)	(39.6)	(42.2)	(84.4)	(97.9)
(180,60)	(186.6, 63.4, 97.5)	(181.2)	(190.2)	(62.2)	(65.7)	(92.7)	(100.3)
(180,80)	(188.4, 74.2, 229.4)	(186.7)	(190.8)	(70.3)	(81.9)	(98.3)	(295.3)
(180,100)	(183.8, 101.9, 146.4)	(181.7)	(186.9)	(99.9)	(104.8)	(89.3)	(314.0)
(180,120)	(176.9, 125.5, 172.2)	(169.3)	(183.3)	(123.9)	(127.1)	(96.1)	(317.3)

Table 6.9: Table with grids for experiment with five tags. Tag 24622

Position	Amount of data points updates	Average deviation
(180, 0)	7	31.70 cm
(180, 20)	6	9.23 cm
(180, 40)	3	4.09 cm
(180, 60)	5	7.61 cm
(180, 80)	3	11.28 cm
(180, 100)	4	4.38 cm
(180, 120)	3	7.90 cm

Table 6.10: Amount of data points with 5 seconds and the average deviation for tag 24622

Tag 26467

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(140,0)	(142.0, -9.4, 229.3)	(132.7)	(147.1)	(-14.2)	(0.0)	(102.3)	(295.7)
(140,20)	(123.0, 3.7, 76.8)	(123.0)	(123.0)	(3.7)	(3.7)	(76.8)	(76.8)
(140,40)	(139.9, 26.8, 94.1)	(136.5)	(142.5)	(24.1)	(30.6)	(92.4)	(95.7)
(140,60)	(123.4, 65.0, 96.0)	(120.8)	(126.0)	(59.7)	(70.4)	(93.4)	(98.7)
(140,80)	(133.1, 70.4, 83.5)	(124.9)	(140.7)	(67.6)	(75.8)	(70.9)	(94.2)
(140,100)	(139.2, 90.6, 88.6)	(138.7)	(139.9)	(85.4)	(99.5)	(77.5)	(94.7)
(140,120)	(137.9, 110.4, 88.8)	(136.3)	(140.1)	(107.6)	(113.2)	(86.6)	(91.2)

Table 6.11: Table with grids for experiment with five tags. Tag 26467

Position	Amount of data points updates	Average deviation
(140, 0)	3	12.32 cm
(140, 20)	1	23.55 cm
(140, 40)	3	13.44 cm
(140, 60)	2	17.91 cm
(140, 80)	4	13.30 cm
(140, 100)	4	9.55 cm
(140, 120)	4	9.98 cm

Table 6.12: Amount of data points with 5 seconds intervals and the average deviation for tag 26467**Tag 26895**

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(160,0)	(174.4, -13.1, 249.4)	(153.6)	(181.6)	(-29.1)	(8.9)	(83.1)	(319.9)
(160,20)	(169.6, 20.8, 87.3)	(165.1)	(177.6)	(11.2)	(29.5)	(77.5)	(98.0)
(160,40)	(164.4, 35.4, 130.4)	(154.3)	(172.9)	(26.8)	(45.8)	(84.4)	(313.3)
(160,60)	(168.2, 62.4, 131.8)	(161.8)	(180.5)	(55.3)	(69.6)	(91.8)	(297.5)
(160,80)	(163.7, 86.9, 96.3)	(158.9)	(166.8)	(79.0)	(96.2)	(90.5)	(105.9)
(160,100)	(166.8, 95.3, 222.9)	(155.1)	(177.5)	(79.1)	(117.9)	(68.7)	(323.2)
(160,120)	(162.1, 133.1, 153.3)	(155.1)	(166.1)	(124.5)	(140.0)	(74.3)	(315.9)

Table 6.13: Table with grids for experiment with five tags. Tag 26895

Position	Amount of data points updates	Average deviation
(160, 0)	11	22.05 cm
(160, 20)	11	10.74 cm
(160, 40)	11	9.31 cm
(160, 60)	11	9.51 cm
(160, 80)	10	8.44 cm
(160, 100)	10	12.75 cm
(160, 120)	10	13.71 cm

Table 6.14: Amount of data points with 5 seconds intervals and the average deviation for tag 26467**Tag 26901**

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(220, 0)	(229.9, 4.9, 233.6)	(214.6)	(246.5)	(-6.9)	(15.7)	(72.4)	(318.3)
(220, 20)	(227.5, 14.8, 151.4)	(221.7)	(233.4)	(3.0)	(30.8)	(87.7)	(296.1)
(220, 40)	(228.6, 42.4, 111.9)	(221.6)	(237.5)	(31.9)	(58.0)	(90.6)	(286.5)
(220, 60)	(227.5, 61.0, 132.9)	(220.5)	(241.1)	(43.7)	(75.2)	(87.1)	(293.6)
(220, 80)	(222.8, 77.6, 86.2)	(218.4)	(235.3)	(60.2)	(90.3)	(75.3)	(91.9)
(220, 100)	(220.2, 92.2, 165.6)	(180.5)	(228.8)	(51.3)	(115.2)	(22.4)	(342.8)
(220, 120)	(222.4, 108.4, 105.6)	(215.6)	(228.7)	(97.7)	(115.5)	(73.9)	(319.7)

Table 6.15: Table with grids for experiment with five tags. Tag 26901

Position	Amount of data points updates	Average deviation
(220, 0)	11	14.79 cm
(220, 20)	11	20.99 cm
(220, 40)	11	14.71 cm
(220, 60)	11	15.61 cm
(220, 80)	10	13.90 cm
(220, 100)	10	8.97 cm
(220, 120)	10	13.54 cm

Table 6.16: Amount of data points with 5 seconds intervals and the average deviation for tag 26467

Tag 27001

Actual grid	Average grid (x,y,z)	x min	x max	y min	y max	z min	z max
(200, 0)	(198.3, 0.3, 133.9)	(194.0)	(204.5)	(-5.8)	(3.8)	(92.7)	(289.2)
(200, 20)	(208.3, 11.7, 181.5)	(201.1)	(220.9)	(-1.7)	(24.2)	(102.7)	(294.6)
(200, 40)	(201.8, 31.8, 113.4)	(195.5)	(206.7)	(16.9)	(42.0)	(108.2)	(118.5)
(200, 60)	(205.6, 53.0, 140.8)	(200.3)	(213.5)	(37.7)	(60.0)	(112.4)	(272.8)
(200, 80)	(203.6, 71.5, 116.4)	(200.1)	(207.9)	(57.4)	(77.4)	(112.5)	(118.9)
(200, 100)	(201.6, 93.2, 112.0)	(197.5)	(205.8)	(90.1)	(96.8)	(109.7)	(115.5)
(200, 120)	(204.6, 110.8, 105.9)	(200.7)	(212.4)	(101.0)	(119.0)	(95.0)	(116.6)

Table 6.17: Table with grids for experiment with five tags. Tag 27001

Position	Amount of data points	Average deviation
(200, 0)	6	7.32 cm
(200, 20)	7	18.34 cm
(200, 40)	7	9.69 cm
(200, 60)	9	8.10 cm
(200, 80)	7	6.21 cm
(200, 100)	3	12.44 cm
(200, 120)	4	14.03 cm

Table 6.18: Amount of data points with 5 seconds intervals and the average deviation for tag 27001

6.3 Test plan

6.3.1 Documentation

All functions should be documented. This is done by writing comments in the code following the standards of the language.

For C#:

<https://docs.microsoft.com/en-us/dotnet/csharp/language-reference/language-specific-documentation-comments>

For Python:

<https://www.geeksforgeeks.org/python-docstrings/>

These comments should explain what the parameters of the function are and a short summary of what the function does.

If the function that is being implemented deals with more complex concepts, additional comments should be written in the code to explain what is happening in the function.

6.3.2 Testing

What to test

All code that can influence the user's experience should be tested. This means that every function that would influence the program negatively if it had an incorrect calculation should be tested to try to avoid this.

All public functions should always be tested.

How to test

- Unit tests
 - Functions are tested in isolation.
 - If the function needs data from outside the function it should be manually generated using mocks.
 - More information about unit testing
- Integration tests
 - Testing if different modules integrate correctly.
 - An example of this is when we are fetching data from the python server into the C# client.
 - More information about integration testing

Unit testing in unity

There are two ways of running unit tests in Unity: play mode and editor mode. Play mode requires the framework to load the scenes before testing code while editor mode can test the scripts in the project without loading the scenes. When the scenes are loaded in play mode, the objects are instantiated with the predefined data in the scene.

Since this makes it more difficult to control the state of the program, the focus will be on testing the functions isolated in editor mode. This has the consequence that some code that utilizes the Unity API and functions on presentational objects in Unity can not be tested without using the play mode.

To make it easier to unit test your code you should try to separate logic that depends on Unity features and logic that does not into different functions.

How much to test

100% code coverage is not necessary, but if there are no unit tests for a part of the system it is not enough. The focus should be on the quality of tests rather than the number of tests.

All functionality in a function that can affect the output of the function should be tested. If a bug is fixed and it took more than a couple of minutes to fix it, unit tests should be created for this bug to prevent it from being reintroduced later on.

Bibliography

- [1] Ivan Aaen. *Essence. Problem Based Digital Innovation*. 2020.
- [2] *Creator System Performance*. URL: <https://www.pozyx.io/products-and-services/creator/system-performance>.
- [3] Finjan Cybersecurity. URL: <https://blog.finjan.com/unicast-broadcast-multicast-data-transmissions/> (visited on 03/01/2020).
- [4] Bill Fenner, Andrew M. Rudoff, and Richards W. Stevens. *UNIX Network Programming, Volume 1. The Sockets Networking API*. 3. edition. Addison-Wesley Professional, 2003.
- [5] Exit Games. URL: <https://www.photonengine.com/en-US/PUN> (visited on 02/14/2019).
- [6] Ilya Grigorik. *High Performance Browser Networking. What Every Developer Should Know About Networking And Web Performance*. 1. edition. O'Reilly Media, 2013.
- [7] *Home: Pozyx*. URL: <https://www.pozyx.io/>.
- [8] *Home: Pozyx*. URL: <https://www.taitradioacademy.com/courses/basic-radio-awareness/>.
- [9] *How does ultra-wideband work?* URL: <https://www.pozyx.io/technology/how-does-uwb-work>.
- [10] infsoft. URL: <https://www.infsoft.com/technology/positioning-technologies/wi-fi> (visited on 03/02/2020).
- [11] infsoft. URL: <https://www.infsoft.com/blog/indoor-navigation-indoor-positioning-using-bluetooth> (visited on 03/02/2020).
- [12] infsoft. URL: <https://www.infsoft.com/technology/positioning-technologies/rfid> (visited on 03/02/2020).
- [13] *IPv4 Multicast Address Space Registry*. URL: <https://www.iana.org/assignments/multicast-addresses/multicast-addresses.xhtml#multicast-addresses-1>.

- [14] I. Scott MacKenzie and Colin Ware. “Lag as a Determinant of Human Performance in Interactive Systems”. In: *Proceedings of the INTERACT 93 and CHI 93 Conference on Human Factors in Computing Systems*. CHI 93. Amsterdam, The Netherlands: Association for Computing Machinery, 1993, pp. 488-493. ISBN: 0897915755. DOI: 10.1145/169059.169431. URL: <https://doi.org/10.1145/169059.169431>.
- [15] Curtiss Murphy. *Believable Dead Reckoning for Networked Games*. URL: https://www.researchgate.net/publication/293809946_Believable_Dead_Reckoning_for_Networked_Games (visited on 04/03/2020).
- [16] US National Coordination Office for Space-based Positioning Navigation and Timing. URL: <https://www.gps.gov/systems/gps/performance/accuracy/> (visited on 03/02/2020).
- [17] *Positioning protocols explained*. URL: <https://www.pozyx.io/technology/positioning-protocols-explained>.
- [18] Unity Technologies. URL: <https://unity.com/how-to/programming-unity> (visited on 02/13/2019).
- [19] Unity Technologies. URL: <https://unity3d.com/unity/features/multiplatform> (visited on 02/13/2019).
- [20] Unity Technologies. URL: <https://docs.unity3d.com/Manual/UNet.html> (visited on 02/14/2019).
- [21] Unity Technologies. *Game engines - how do they work?* URL: <https://unity3d.com/what-is-a-game-engine> (visited on 02/13/2019).
- [22] *Ultra-wideband and obstacles*. URL: <https://www.pozyx.io/technology/uwb-and-obstacles>.
- [23] Thomas Waltemate et al. “The impact of latency on perceptual judgments and motor performance in closed-loop interaction in virtual reality”. eng. In: *Proceedings of the 22nd ACM Conference on virtual reality software and technology*. Vol. 02-04-. VRST '16. ACM, 2016, pp. 27–35. ISBN: 9781450344913.
- [24] *Where to place the anchors*. URL: <https://www.pozyx.io/technology/where-to-place-the-anchors>.
- [25] ZeroMQ. URL: <https://zeromq.org/get-started/> (visited on 02/18/2019).
- [26] ZeroMQ. URL: <https://zguide.zeromq.org/page:all> (visited on 02/18/2019).

List of Figures

1.1	An illustration of the playing field	1
2.1	The board for tasks.	7
2.2	Prototype of hosting interface	8
2.3	Prototype of game menu	8
2.4	Prototype of screen where a user has connected to the host	9
2.5	Prototype of in-game screen	9
2.6	The different components of the game	10
2.7	An example of trilateration using three anchors	11
2.8	An illustration of the OSI-layer, and the corresponding technology for each layer.	18
2.9	An illustration of the client-server architecture with multiple clients .	19
2.10	An illustration of the client-server architecture with UDP	22
2.11	An illustration of a publisher-subscriber architecture with UDP	22
2.12	The setup of the experiment with the anchors and the height at which they were placed in the corners. The hexadecimal number is the anchor, and above that is the height at which the anchors were placed. .	26
2.13	The whiteboard with the drawn positions.	26
2.14	A box plot showing the dispersion of measured positions with 1 tag . .	27
2.15	A box plot showing the dispersion of measured positions with 3 tags .	29
2.16	A box plot showing the dispersion of measured positions with 5 tags .	30
3.1	A deployment diagram for the system.	36
3.2	Prototype of the game menu when players can choose a game to join.	38
3.3	Prototype of a lobby menu. Available games are shown as well as the number of players. Players can choose one to join.	38
3.4	The first implementation of VR	41
4.1	Linear player movement	46
4.2	A new update about the player's state is received from the server . . .	46
4.3	Shows the the curve for the dead reckoned position Q_t when making use of projective velocity blending	48

List of Tables

2.1	A comparison of the pros and cons of the possible solutions	17
2.2	Format for field anchor position	24
2.3	Format for player positions.	24
2.4	Format for the goal positions.	24
2.5	Format for goals scored.	25
6.1	Table with grids for experiment with 1 tag. Tag: 26895	52
6.2	Average deviation for each position with tag 26895	53
6.3	Table with grids for experiment with three tags. Tag: 26467	53
6.4	Average deviation for tag 26467	54
6.5	Table with grids for experiment with three tags. Tag: 26895	54
6.6	Average deviation for tag 26895	54
6.7	Table with grids for experiment with three tags. Tag: 24622	54
6.8	Average deviation for tag 24622	55
6.9	Table with grids for experiment with five tags. Tag 24622	55
6.10	Amount of data points with 5 seconds and the average deviation for tag 24622	55
6.11	Table with grids for experiment with five tags. Tag 26467	56
6.12	Amount of data points with 5 seconds intervals and the average deviation for tag 26467	56
6.13	Table with grids for experiment with five tags. Tag 26895	56
6.14	Amount of data points with 5 seconds intervals and the average deviation for tag 26467	57
6.15	Table with grids for experiment with five tags. Tag 26901	57
6.16	Amount of data points with 5 seconds intervals and the average deviation for tag 26467	57
6.17	Table with grids for experiment with five tags. Tag 27001	58
6.18	Amount of data points with 5 seconds intervals and the average deviation for tag 27001	58

Listings

3.1	Updating ball position	39
3.2	Receiving data from host	39